



May 2025. An AI agent read one GitHub issue and published its owner's private code. Every step it took was allowed.



Nothing was hacked. The agent was obeyed. Per-tool permissions can't see this: **the danger is the sequence.**

GitHub MCP, May 2025 (Invariant Labs): the issue hid instructions for the agent. Same pattern, Supabase MCP, Jul 2025: a support ticket made the agent read private tables and write them back into the ticket.



“The session read text written by outsiders and also holds private data, so sending data out could leak it.”

Every role gets its own MCP server from one policy file; every call is checked locally; once a session has read untrusted text, **private** → **public** is blocked.

Scenario act 3 (a3-attack) k84155fd2Active

Agent	Role	Started	Duration	Steps	Blocked	Masked	State
Support assistant	role-2	03:40:11	<1 s	3	1	1	⚠ Untrusted 🔒 Holds private data

Steps

List

Timeline

1	Read an issue	71 ms
2	Read a file from a repository	🔒 Masked 1.4 ms
3	Open a pull request	🚫 Blocked 0.5 ms

Open a pull request: was blocked.

03:40:12 🚫 Blocked Data-flow rule: blocked 0.5 ms

Overview


Input/Output

Checks

Sent to server

What happened	Open a pull request: was blocked.
Check	Data-flow rule: blocked
Effect on the session	No change to the session.
Why	The session read text written by outsiders and also holds private data, so sending data out could leak it.
What you can do	Start a new session (new key) for this task, without reading untrusted text.

One policy file, **one virtual MCP per role**




Hybrid checks, **taint** as the backstop

Layer	What it catches	Held-out result
PII & secrets, deterministic	IBAN, PESEL, cards (checksums: Luhn, mod-97); AWS keys	PII recall 0.967
Injection regex + normalisation	zero-width chars, base64 / hex / URL encoding undone first	obfuscated: 1.00 (0.67 without)
DeBERTa classifier, local ONNX	paraphrased prompt injection	recall 0.837 , FPR 0.009
Signatures, signed feed	historical attacks: <code>pickle</code> , <code>torch.load</code> , <code>trust_remote_code</code>	hot-loaded, no restart
Session taint	what text checks miss: untrusted + private can't reach a public sink	both 2025 attacks blocked

All checks together: recall **0.891**, FPR **0.048**. Posture score **0.931**.

Every attack, traced to its session



Open a pull request: was blocked.
03:40:12 Blocked Data-flow rule: blocked 0.5 ms

Overview Checks **Fingerprint**

🔒 The text stays on the peer's laptop. The server keeps the decision and this fingerprint only.

Fingerprint (SHA-256)	56992091beed1b20fb3d6f620b2af6ff60efe07ee7b6c6865ce88b842768962e	Copy
Decision	Blocked	
Check	Data-flow rule: blocked	
Tool	Open a pull request	
Peer	Unenrolled	
Session state	Untrusted + Holds private data	
Time	03:40:12	
Trace ID	t_0cd265067e37	

Click an attack → its session, step by step, cause in plain words. **Text stays on the laptop**: the server keeps the decision and a SHA-256 fingerprint.

Edit access in the console, **the next call obeys**

Review access changes

Intern bot role-1

- Intern bot may now open a pull request (write, public destination)

⚠ Intern bot can now read outside text and private data and send to a public destination. The data-flow rule still guards this: when a session has read outside text and holds private data, a call to a public destination is blocked.

Saving backs up the policy file, writes the new access and applies it from the next call. Comments in the file are not kept; the backup has them.

Back to editingSave 1 change

- Matrix: role × MCP server × tool.
- Review shows the diff and warns about risky data flows.
- Hot reload on the next call, no restart.
- A broken edit is **rejected**; the old policy keeps enforcing.
- Profiles: strict / balanced / lenient.

Enroll once, a key per agent launch

Laptop	Roles	Seen	Sessions	Blocked
andrey-thinkpad Andrey Mazurchuk	Support assistant	● Online	4	2
kasia-x1 Kasia Lewandowska	Intern bot	● Online	2	2
marta-macbook Marta Kowalska	Support assistant	● Online	2	6
ola-surface Ola Wiśniewska	Intern bot, Support a...	● Online	2	2
piotr-xps Piotr Nowak	Intern bot	● Online	2	2
tomek-mbp Tomasz Zieliński	Support assistant	● Online	6	0

- A laptop enrolls once; the admin sets the roles it may run.
- Each agent launch mints its own key.
- Revoke a laptop: every key stops at once.
- Works with **any MCP server**: the Hub fronts it per role.

Judges run it: `uv run tollgate test`

247

fast tests pass

(+7 slow: real model, full eval)

1,411

eval cases: public corpora + own red team,

30% held out, tuned on the rest

0.837

injection recall vs target 0.85:

missed, stated plainly

Injection recall

83.7%

✗ Target $\geq 85.0\%$: missed

eval, held-out 30%

Injection false alarms

0.9%

✓ Target $\leq 5.0\%$: met

harmless text flagged, held-out 30%

False alarms, all checks

4.8%

✓ Target $\leq 5.0\%$: met

harmless cases flagged, held-out 30%

Personal-data recall

96.7%

✓ Target $\geq 95.0\%$: met

eval, held-out 30%

Tier 1 scan p95

2.7 ms

✓ Target ≤ 5.0 ms: met

eval

Classifier p95, short text

87 ms

✗ Target ≤ 80 ms: missed

eval, ≤ 64 tokens

The Hub costs **under a millisecond**

< 1 ms

Hub checks p95
(roles, limits, taint)

2.5 ms

tier 1 p95
(regex, checksums, signatures)

65–90 ms

classifier p95, short text:
borderline vs 80 ms target

- The classifier is **gated**: prompts always get it; long tool results are deferred, so calls stay fast.
- Stateless checks per request; one policy for every role; taint is one small record per key.
- `uv run tollgate perf` reproduces the numbers; the console Self-test shows the last run.

```
git clone https://github.com/andrii-mazurchuk/tollgate && cd tollgate
uv sync
uv run tollgate up --scripted-model # Hub + mock MCPs + offline model; /console, /edge
uv run tollgate test # 247 tests + eval report
```

Then open `/edge/` → Scenario → Run all: both 2025 attacks replay and get blocked.

WHAT'S NEXT

- A live access map in the console: every laptop, agent and tool call on one interactive view.
- Cost budgets in money, not only tokens.
- Memory controls for agents with long-term memory.